



# SISTEMI ZA OBRADU I ANALIZU VELIKE KOLIČINE PODATAKA - Projekat 2

Filip Trajković 1574

# • KORIŠĆENE TEHNOLOGIJE

## Apache Hadoop



## Apache Flink



## Apache Spark



## Cassandra



# Spisak kontejnera

<input type="checkbox"/>	 project2	-	Running (11/1)	  
<input type="checkbox"/>	 producer-1 b215a72cf2ed 	<a href="#">project2-producer</a>	Running	21 seconds ago   
<input type="checkbox"/>	 kafka-1 9f96a8c12b69 	<a href="#">wurstmeister/kafka:2</a>	Running	9091:9091  22 seconds ago   
<input type="checkbox"/>	 taskmanager-1 702bba3d5922 	<a href="#">flink:latest</a>	Running	21 seconds ago   
<input type="checkbox"/>	 spark-worker-1-x 76cf883715e 	<a href="#">bde2020/spark-work</a>	Running	8071:8071  22 seconds ago   
<input type="checkbox"/>	 spark-worker-2-x b7d763b69dbd 	<a href="#">bde2020/spark-work</a>	Running	8072:8071  22 seconds ago   
<input type="checkbox"/>	 zookeeper-1 2ca94cdd4bb1 	<a href="#">wurstmeister/zookeeper:3</a>	Running	2181:2181  22 seconds ago   
<input type="checkbox"/>	 consumer_flink-1 4ce7be3489aa 	<a href="#">project2-consumer_flink</a>	Running	8082:8082  22 seconds ago   
<input type="checkbox"/>	 spark-master-x 6fc41023a9ef 	<a href="#">bde2020/spark-master</a>	Running	7077:7077  8070:8070  23 seconds ago   
<input type="checkbox"/>	 cassandra-1 dc2b81865a3c 	<a href="#">cassandra:4.0</a>	Running	9042:9042  22 seconds ago   
<input type="checkbox"/>	 jobmanager-1 2355fc2c8ed4 	<a href="#">flink:latest</a>	Running	8081:8081  22 seconds ago   
<input type="checkbox"/>	 consumer_spark-1 aa0b85f591c9 	<a href="#">project2-consumer_spark</a>	Running	4040:4040  22 seconds ago   



# Izgled Producer skripte



```
8 producer = KafkaProducer(  
9     bootstrap_servers=[os.environ["KAFKA_HOST"]],  
10    value_serializer=lambda v: json.dumps(v).encode("utf-8"),  
11    api_version=(0, 11),  
12 )  
13  
14 while True:  
15     with open(os.environ["DATA"], "r") as file:  
16         reader = csv.reader(file, delimiter=",")  
17         headers = next(reader)  
18         for row in reader:  
19             value = {headers[i]: row[i] for i in range(len(headers))}  
20             value["lat"] = float(value["lat"])  
21             value["lon"] = float(value["lon"])  
22             value["alt"] = float(value["alt"])  
23             value["user"] = int(value["user"])  
24             value["ts"] = int(time.time())  
25             value["year"] = int(value["year"])  
26             value["month"] = int(value["month"])  
27             value["day"] = int(value["day"])  
28             value["hour"] = int(value["hour"])  
29             value["minute"] = int(value["minute"])  
30             value["second"] = int(value["second"])  
31             print(value)  
32             producer.send(os.environ["KAFKA_TOPIC"], value=value)  
33             time.sleep(float(os.environ["KAFKA_INTERVAL"]))
```

# Izgled Cassandra tabela



```
# cqlsh
Connected to cloudinfra at 127.0.0.1:9042
[cqlsh 6.0.0 | Cassandra 4.0.7 | CQL spec 3.4.5 | Native protocol v5]
Use HELP for help.
cqlsh> DESC keyspaces;

locations_db  system_auth      system_schema  system_views
system        system_distributed system_traces   system_virtual_schema

cqlsh> USE locations_db;
cqlsh:locations_db> DESC tables;

flinkgeolocation  geolocation

cqlsh:locations_db> SELECT * FROM geolocation;

user | end                  | maxvalue | meanvalue          | minvalue | start                | times_data_sent
-----+-----+-----+-----+-----+-----+-----+-----+
 163 | 2023-03-30 22:56:00+0000 | 190.28871 | 190.28871154785156 | 190.28871 | 2023-03-30 22:55:50+0000 | 2
 154 | 2023-02-15 16:53:00+0000 | 16.4042   | 16.404199600219727 | 16.4042   | 2023-02-15 16:52:50+0000 | 8
  21 | 2023-02-15 16:38:00+0000 | 1036.7454  | 1033.464564732143  | 1030.1837  | 2023-02-15 16:37:50+0000 | 7
 111 | 2023-03-30 23:00:55+0000 | 160.76115  | 158.5739288330078  | 154.19948  | 2023-03-30 23:00:45+0000 | 3
   97 | 2023-03-24 17:16:40+0000 | 249.34383  | 242.7821502685547  | 236.22047  | 2023-03-24 17:16:30+0000 | 2
   86 | 2023-02-15 16:38:10+0000 | 187.00787  | 187.00787353515625 | 187.00787  | 2023-02-15 16:38:00+0000 | 2
 118 | 2023-02-15 17:33:30+0000 | 170.60367  | 158.7926483154297  | 141.07611  | 2023-02-15 17:33:20+0000 | 5
 161 | 2023-03-30 22:58:55+0000 | 269.02887  | 260.27996826171875 | 255.90552  | 2023-03-30 22:58:45+0000 | 3
 128 | 2023-03-30 23:02:45+0000 | 147.6378   | 142.71653747558594 | 137.79527  | 2023-03-30 23:02:35+0000 | 2

(9 rows)
cqlsh:locations_db> SELECT * FROM flinkgeolocation;

user | count | maxvalue | meanvalue | minvalue
-----+-----+-----+-----+-----+
 163 | 3.0   | 160.761154855643 | 160.761154855643 | 160.761154855643
 111 | 2.0   | 160.761154855643 | 160.761154855643 | 160.761154855643
 161 | 1.0   | 255.905511811024 | 255.905511811024 | 255.905511811024
 128 | 2.0   | 187.007874015748 | 180.446194225722 | 173.884514435696

(4 rows)
cqlsh:locations_db>
```

# Streaming aplikacija

## Shell skripta za pokretanje kontejnera

```
1  #! /bin/bash
2
3  docker container rm SparkStreamingApp
4
5  docker build --rm -t bde/spark-app .
6
7  docker run --name SparkStreamingApp --net project2_default -p 4042:4042 bde/spark-app
```

Expertise

Expertise

Expertise



# Streaming aplikacija

## Dockerfile

```
1 FROM bde2020/spark-python-template:3.1.2-hadoop3.2
2
3 ENV KAFKA_HOST=kafka:9092
4 ENV KAFKA_TOPIC=locations
5 ENV KAFKA_CONSUMER_GROUP=Spark-Group
6 ENV SPARK_APPLICATION_PYTHON_LOCATION /app/streaming_app.py
7 ENV SPARK_APPLICATION_ARGS "lat lon | 38.0 40.1 115.4 130.6"
8 ENV SPARK_SUBMIT_ARGS --packages org.apache.spark:spark-sql-kafka-0-10_2.12:3.1.2,com.datastax.spark:spark-cassandra-connector_2.12:3.2.0
```



# Streaming aplikacija

## Setup Spark i Kafka konekcija

```
73  if __name__ == '__main__':
74      if len(sys.argv) < 2:
75          print("Usage: main.py <input folder> ")
76          exit(-1)
77
78      inputValuesDictionary = readInputValues()
79      print(inputValuesDictionary)
80
81
82      conf = SparkConf()
83      #conf.setMaster("spark://spark-master-x:7077")
84      conf.setMaster("local")
85      conf.set("spark.driver.memory", "4g")
86      conf.set("spark.cassandra.connection.host", "cassandra")
87      conf.set("spark.cassandra.connection.port", 9042)
88
89      spark = SparkSession.builder.config(conf=conf).appName("AverageLatitude").getOrCreate()
90
91      spark.sparkContext.setLogLevel("ERROR")
92      Expertise
93      df = (
94          spark.readStream.format("kafka")
95          .option("kafka.bootstrap.servers", "kafka:9092")
96          .option("subscribe", "locations")
97          .option("startingOffsets", "latest")
98          .load()
99      )
100
```

Expertise →

# Streaming aplikacija

```
103     parsed_valuesWithTimestamp = df.select(  
104         "timestamp", from_json(col("value").cast("string"), schema).alias("parsed_values"))  
105     )  
  
106  
107     parsed_valuesWithTimestamp.printSchema()  
108     #parsed_valuesWithTimestamp = parsed_valuesWithTimestamp.where(col("parsed_values.lat") > inputValuesDictionary["lat"][0]).where(col("parsed_values.lat")  
109     parsed_valuesWithTimestamp = findDfBasedOnInputValues(inputValuesDictionary, parsed_valuesWithTimestamp)  
110  
111     timestampXaltitudes = parsed_valuesWithTimestamp.selectExpr("timestamp", "parsed_values.alt AS alt", "parsed_values.user AS user")  
112  
113     parsed_values = parsed_valuesWithTimestamp.selectExpr("parsed_values.*")  
114     #parsed_values.printSchema()  
115     timestamps = parsed_values.selectExpr("ts AS ts")  
116     latitudes = parsed_values.selectExpr("lat AS lat")  
117     longitudes = parsed_values.selectExpr("lon AS lon")  
118     altitudes = parsed_values.selectExpr("alt AS alt")  
119     labels = parsed_values.selectExpr("label AS label")  
120     users = parsed_values.selectExpr("user AS user")  
121  
122  
123  
124  
125     #timestampXaltitudes.printSchema()  
126     AltitudeValuesAvgWindows = timestampXaltitudes.groupBy(window(timestampXaltitudes.timestamp, "20 seconds", "10 seconds"), timestampXaltitudes.user).agg(  
127         mean("alt").alias("meanvalue"),  
128         max("alt").alias("maxvalue"),  
129         min("alt").alias("minvalue"),  
130         count("user").alias("times_data_sent"))  
131     )  
132     Expertise  
133  
134     #AltitudeValuesAvgWindows.printSchema()  
135     AltitudeValuesAvgWindowsForDB = AltitudeValuesAvgWindows.selectExpr("window.start as start",  
136                             "window.end as end",  
137                             "user",  
138                             "meanvalue",  
139                             "maxvalue",  
140                             "minvalue",  
141                             "times_data_sent" )  
142     #AltitudeValuesAvgWindowsForDB.printSchema()  
143  
144
```

Expertise

Konverzija  
ulaznih  
vrednosti



# Streaming aplikacija

Upis rezultata  
obrade  
na konzolu  
i u  
Cassandra  
bazu  
podataka

```
145  ↴    query = (  
146  ↴        AltitudeValuesAvgWindows.writeStream.outputMode("update")  
147  
148        .queryName("average_latitude")  
149        .format("console") # format("org.apache.spark.sql.cassandra").options(**load_options)  
150        .trigger(processingTime="1 seconds")  
151        .option("truncate", "false")  
152        .start()  
153    )  
154  
155 #####  
156 ##### UPIS U CASSANDRA BAZU PODATAKA #####  
157 #####  
158  
159  ↴    def writeToCassandra(writeDF, _):  
160  ↴        writeDF.write \  
161            .format("org.apache.spark.sql.cassandra")\  
162            .mode('append')\  
163            .options(table="geolocation", keyspace="locations_db")\  
164        Expertise.save()  
165  
166  ↴    AltitudeValuesAvgWindowsForDB.writeStream \  
167        .option("spark.cassandra.connection.host","cassandra:9042")\  
168        .foreachBatch(writeToCassandra) \  
169        .outputMode("update") \  
170        .start() \  
171        .awaitTermination()  
172
```

Expertise



# Streaming aplikacija

# Prikaz batch obrade u konzoli

Batch: 6

		user meanvalue	maxvalue	minvalue	times_data_sent
+-----+-----+-----+-----+-----+	window				
+-----+-----+-----+-----+-----+	{2023-04-04 17:33:20, 2023-04-04 17:33:30} 128	229.65879821777344	232.93964	226.37796	2
+-----+-----+-----+-----+-----+	{2023-04-04 17:33:25, 2023-04-04 17:33:35} 128	229.65879821777344	232.93964	226.37796	2
+-----+-----+-----+-----+-----+					

Batch: 7

		user meanvalue	maxvalue	minvalue	times_data_sent
+-----+-----+-----+-----+-----+	window				
+-----+-----+-----+-----+-----+	{2023-04-04 17:33:35, 2023-04-04 17:33:45} 128	254.2650909423828	282.15222	226.37796	2
+-----+-----+-----+-----+-----+	{2023-04-04 17:33:30, 2023-04-04 17:33:40} 128	221.45669555664062	226.37796	216.53543	2
+-----+-----+-----+-----+-----+	{2023-04-04 17:33:25, 2023-04-04 17:33:35} 128	225.28434244791666	232.93964	216.53543	3
+-----+-----+-----+-----+-----+	{2023-04-04 17:33:40, 2023-04-04 17:33:50} 128	282.1522216796875	282.15222	282.15222	1
+-----+-----+-----+-----+-----+					

Batch: 8

		user meanvalue	maxvalue	minvalue	times_data_sent
+-----+-----+-----+-----+-----+	window				
+-----+-----+-----+-----+-----+	{2023-04-04 17:33:45, 2023-04-04 17:33:55} 128	257.5459289550781	259.18634	255.90552	2
+-----+-----+-----+-----+-----+	{2023-04-04 17:33:40, 2023-04-04 17:33:50} 128	265.74802652994794	282.15222	255.90552	3
+-----+-----+-----+-----+-----+					

Batch: 9

		user meanvalue	maxvalue	minvalue	times_data_sent
+-----+-----+-----+-----+-----+	window				
+-----+-----+-----+-----+-----+	{2023-04-04 17:33:45, 2023-04-04 17:33:55} 128	252.62466939290366	259.18634	242.78215	3
+-----+-----+-----+-----+-----+	{2023-04-04 17:33:50, 2023-04-04 17:34:00} 128	264.1076126098633	285.43307	242.78215	2
+-----+-----+-----+-----+-----+	{2023-04-04 17:33:55, 2023-04-04 17:34:05} 128	285.4330749511719	285.43307	285.43307	1
+-----+-----+-----+-----+-----+					

# Flink Streaming aplikacija

## Eksterni "Dependency" u Flink projektu

```
85      </dependency>
86
87      <dependency>
88          <groupId>org.slf4j</groupId>
89          <artifactId>slf4j-api</artifactId>
90          <version>1.7.5</version>
91      </dependency>
92      <dependency>
93          <groupId>org.slf4j</groupId>
94          <artifactId>slf4j-log4j12</artifactId>
95          <version>1.7.5</version>
96      </dependency>
97
98      <dependency>
99          <groupId>junit</groupId>
100         <artifactId>junit</artifactId>
101        <version>4.11</version>
102        <scope>test</scope>
103    </dependency>
104    <dependency>
105        <groupId>org.projectlombok</groupId>
106        <artifactId>lombok</artifactId>
107        <version>RELEASE</version>
108        <scope>compile</scope>
109    </dependency>
110
111  </dependencies>
```

Expertise

Expertise

```
54      <dependencies>
55
56          <dependency>
57              <groupId>org.apache.flink</groupId>
58              <artifactId>flink-scala_2.12</artifactId>
59              <version>${flink.version}</version>
60          </dependency>
61          <dependency>
62              <groupId>org.apache.flink</groupId>
63              <artifactId>flink-streaming-scala_2.12</artifactId>
64              <version>${flink.version}</version>
65          </dependency>
66          <dependency>
67              <groupId>org.apache.flink</groupId>
68              <artifactId>flink-connector-kafka</artifactId>
69              <version>${flink.version}</version>
70          </dependency>
71          <dependency>
72              <groupId>org.apache.flink</groupId>
73              <artifactId>flink-clients</artifactId>
74              <version>${flink.version}</version>
75          </dependency>
76          <dependency>
77              <groupId>org.apache.flink</groupId>
78              <artifactId>flink-streaming-java</artifactId>
79              <version>${flink.version}</version>
80          </dependency>
81          <dependency>
82              <groupId>org.apache.flink</groupId>
83              <artifactId>flink-connector-cassandra_2.12</artifactId>
84              <version>${flink.version}</version>
85          </dependency>
```



# Flink Streaming aplikacija

## Postavljanje stream-a

```
2 usages ▲ MrFICAX*
57 ► public class DataStreamJob {
58
no usages ▲ MrFICAX*
59 ►   public static void main(String[] args) throws Exception {
60
61     StreamExecutionEnvironment environment = StreamExecutionEnvironment.getExecutionEnvironment();
62     String inputTopic = "locations";
63     String server = "kafka:9092";
64
65     DataStream<String> dataStream = StreamConsumer(inputTopic, server, environment);
66
new *
67     DataStream<Location> locationStream = ConvertStreamFromJsonToLocationType(dataStream).filter(new FilterFunction<Location>() {
68         new *
69         @Override
70         public boolean filter(Location location) throws Exception {
71             return (location.getLongitude() > 115.2 && location.getLongitude() < 117.5);
72         });
73     //locationStream.print();
74
75     SingleOutputStreamOperator<WindowedStream<Location, String, TimeWindow>> windowedStream = locationStream
76         .keyBy(Location::getUser) KeyedStream<Location, String>
77         .window(SlidingProcessingTimeWindows.of(Time.seconds(10), Time.seconds(5))) WindowedStream<Location, String, TimeWindow>
78         .aggregate(new AverageAggregate());
79
80     SingleOutputStreamOperator<WindowedStream<Location, String, TimeWindow>> windowedStream2 = locationStream
81         .keyBy(Location::getUser) KeyedStream<Location, String>
82         .window(SlidingProcessingTimeWindows.of(Time.seconds(10), Time.seconds(5))) WindowedStream<Location, String, TimeWindow>
83         .process(new MyProcessWindowFunction());
84         //.aggregate(new AverageAggregate(), new MyProcessWindowFunction());
85
86     //windowedStream.print();
87     //windowedStream2.print();
88
89     CassandraService cassandraService = new CassandraService();
90     cassandraService.sinkToCassandraDB(windowedStream);
91
92     environment.execute();
93 }
```



# Flink Streaming aplikacija

## Kreiranje Consumer-a sa Kafka Topic-a

```
1 usage  ▲ MrFICAX
7 @      public static DataStream<String> StreamConsumer(String inputTopic, String server, StreamExecutionEnvironment environment) throws Exception {
8     FlinkKafkaConsumer<String> flinkKafkaConsumer = createStringConsumerForTopic(inputTopic, server);
9     DataStream<String> stringInputStream = environment.addSource(flinkKafkaConsumer);
10
11
12     ▲ MrFICAX
13     return stringInputStream.map(new MapFunction<String, String>() {
14         no usages
15         private static final long serialVersionUID = -999736771747691234L;
16
17         ▲ MrFICAX
18         @Override
19         public String map(String value) throws Exception {
20             return value;
21         }
22     });
23 }
24
25
26     ▲ MrFICAX
27     public static FlinkKafkaConsumer<String> createStringConsumerForTopic(
28         String topic, String kafkaAddress) {
29
30         Properties props = new Properties();
31         props.setProperty("bootstrap.servers", kafkaAddress);
32         //props.setProperty("group.id", kafkaGroup);
33         FlinkKafkaConsumer<String> consumer = new FlinkKafkaConsumer<>(
34             topic, new SimpleStringSchema(), props);
35
36         return consumer;
37     }
38 }
```



# Flink Streaming aplikacija

## Location klasa

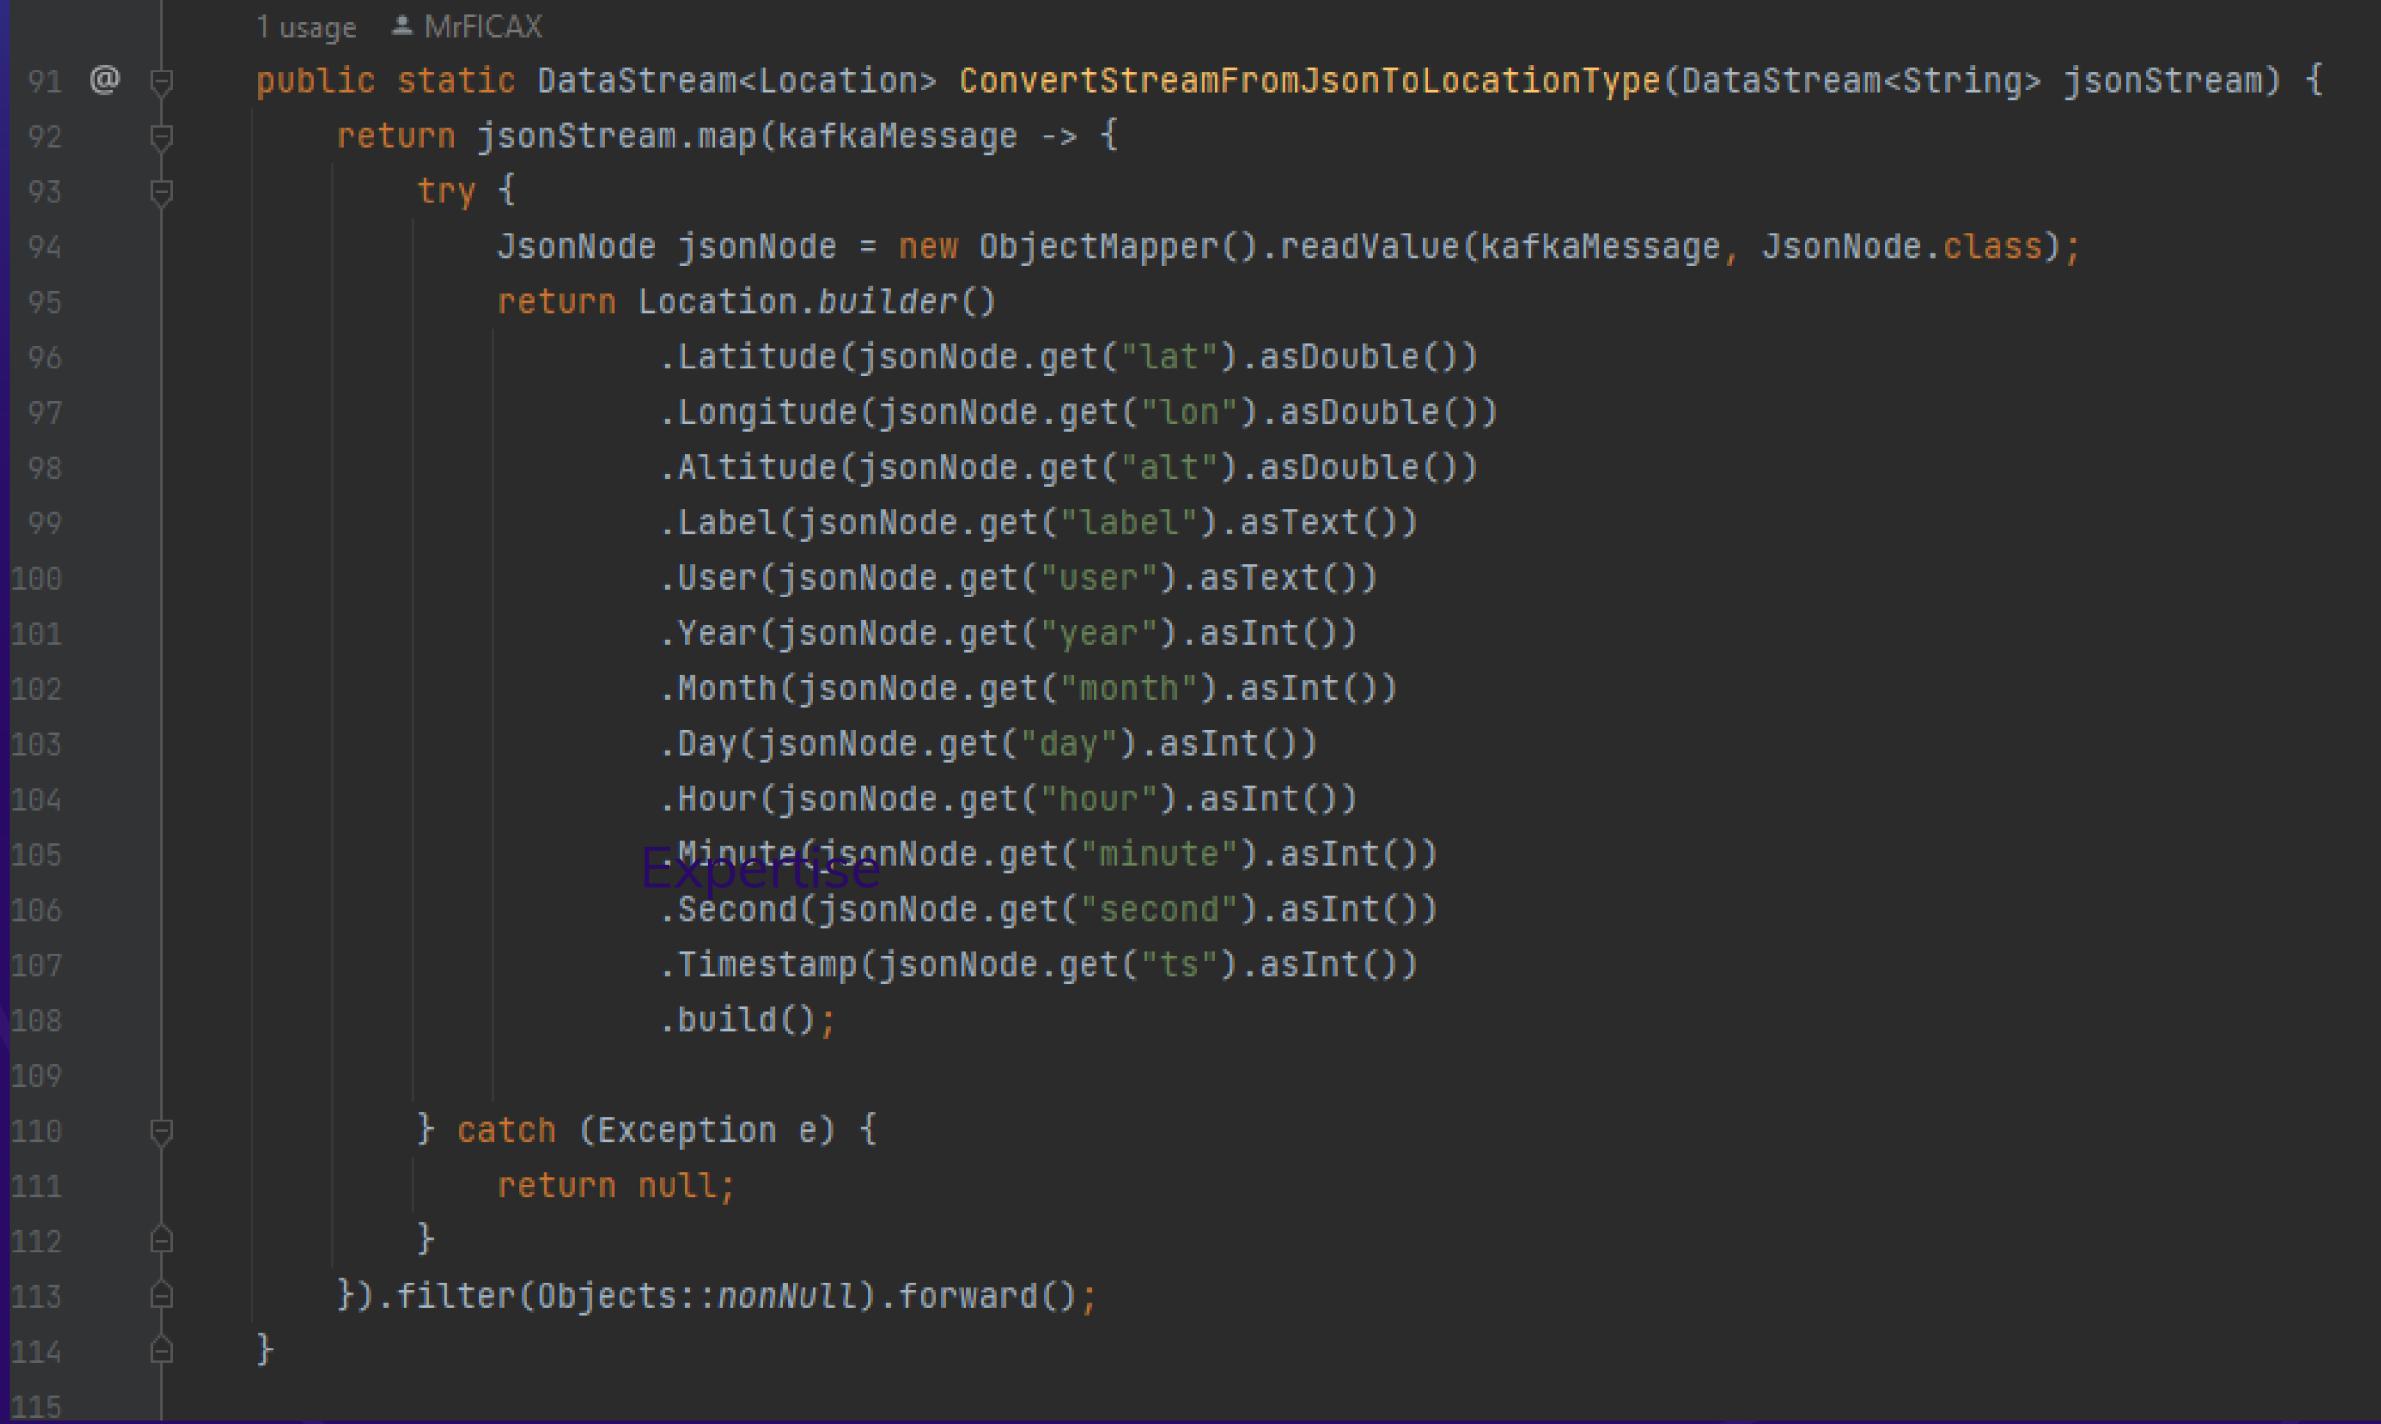
```
4  @Builder
5  @Getter
6  @Setter
7  @ToString
8  @EqualsAndHashCode
9  @AllArgsConstructor
10 @NoArgsConstructor(force = true)
11 public class Location {
12     no usages
13     final Double Latitude;
14     no usages
15     final Double Longitude;
16     no usages
17     final Double Altitude;
18     no usages
19     final String Label;
20     no usages
21     final String User;
22     no usages
23     final Integer Year;
24     no usages
25     final Integer Month;
26     no usages
27     final Integer Day;
28     no usages
29     final Integer Hour;
30     no usages
31     final Integer Minute;
32     no usages
33     final Integer Second;
34     no usages
35     final Integer Timestamp;
36 }
```



# Flink Streaming aplikacija

## Konverzija Stream-a uz Json tipa u tip Location

```
1 usage  ▲ MrFICAX
91  @  ↴
92    ↴
93    ↴
94      ↴
95        ↴
96          ↴
97            ↴
98              ↴
99                ↴
100               ↴
101                 ↴
102                   ↴
103                     ↴
104                       ↴
105                         ↴
106                           ↴
107                             ↴
108                               ↴
109
110     ↴
111       ↴
112         ↴
113           ↴
114             ↴
115               ↴
```



```
public static DataStream<Location> ConvertStreamFromJsonToLocationType(DataStream<String> jsonStream) {
    return jsonStream.map(kafkaMessage -> {
        try {
            JsonNode jsonNode = new ObjectMapper().readValue(kafkaMessage, JsonNode.class);
            return Location.builder()
                .Latitude(jsonNode.get("lat").asDouble())
                .Longitude(jsonNode.get("lon").asDouble())
                .Altitude(jsonNode.get("alt").asDouble())
                .Label(jsonNode.get("label").asText())
                .User(jsonNode.get("user").asText())
                .Year(jsonNode.get("year").asInt())
                .Month(jsonNode.get("month").asInt())
                .Day(jsonNode.get("day").asInt())
                .Hour(jsonNode.get("hour").asInt())
                .Minute(jsonNode.get("minute").asInt())
                .Expertise
                .Second(jsonNode.get("second").asInt())
                .Timestamp(jsonNode.get("ts").asInt())
                .build();
        } catch (Exception e) {
            return null;
        }
    }).filter(Objects::nonNull).forward();
}
```



# Flink Streaming aplikacija

## AggregateFunction za obradu svakog prozora

```
1 usage  ± MrFICAX
class AverageAggregate implements AggregateFunction<Location, Tuple5<String, Double, Double, Double, Double>, Tuple5<String, Double, Double, Double, Double>> {
    no usages  ± MrFICAX
    @Override
    public Tuple5<String, Double, Double, Double, Double> createAccumulator() {
        return new Tuple5<>("", 0D, Double.MAX_VALUE, Double.MIN_VALUE, 0D);
    }

    ± MrFICAX
    @Override
    public Tuple5<String, Double, Double, Double, Double> add(Location value, Tuple5<String, Double, Double, Double, Double> accumulator) {
        ●   accumulator.f0 = value.getUser();
        accumulator.f1 += value.getAltitude();
        accumulator.f2 = accumulator.f2 > value.getAltitude() ? value.getAltitude() : accumulator.f2; //MINIMUM
        accumulator.f3 = accumulator.f3 < value.getAltitude() ? value.getAltitude() : accumulator.f3; //MAXIMUM

        return new Tuple5<>(accumulator.f0, accumulator.f1, accumulator.f2, accumulator.f3, accumulator.f4 + 1);
    }

    ± MrFICAX
    @Override
    public Tuple5<String, Double, Double, Double, Double> getResult(Tuple5<String, Double, Double, Double, Double> acc) {
        return new Tuple5<>(acc.f0, acc.f2, acc.f3, acc.f1 / acc.f4, acc.f4);
    }

    ± MrFICAX
    @Override
    public Tuple5<String, Double, Double, Double, Double> merge(Tuple5<String, Double, Double, Double, Double> acc1, Tuple5<String, Double, Double, Double, Double> acc2) {
        return new Tuple5<>(acc1.f0, acc1.f1 + acc2.f1, acc1.f2 < acc2.f2 ? acc1.f2 : acc2.f2, acc1.f2 > acc2.f2 ? acc1.f2 : acc2.f2, acc1.f4 + acc2.f4);
    }
```



# Flink Streaming aplikacija

## Konverzija DataStream-a i upis u Cassandra bazu podataka

```
public final class CassandraService {

    2 usages
    private static final Logger LOGGER = Logger.getLogger(DataStreamJob.class);

    1 usage  ↳ MrFICAX*
    public final void sinkToCassandraDB(final DataStream<Tuple5<String, Double, Double, Double, Double>> sinkFilteredLocationStream) throws Exception {

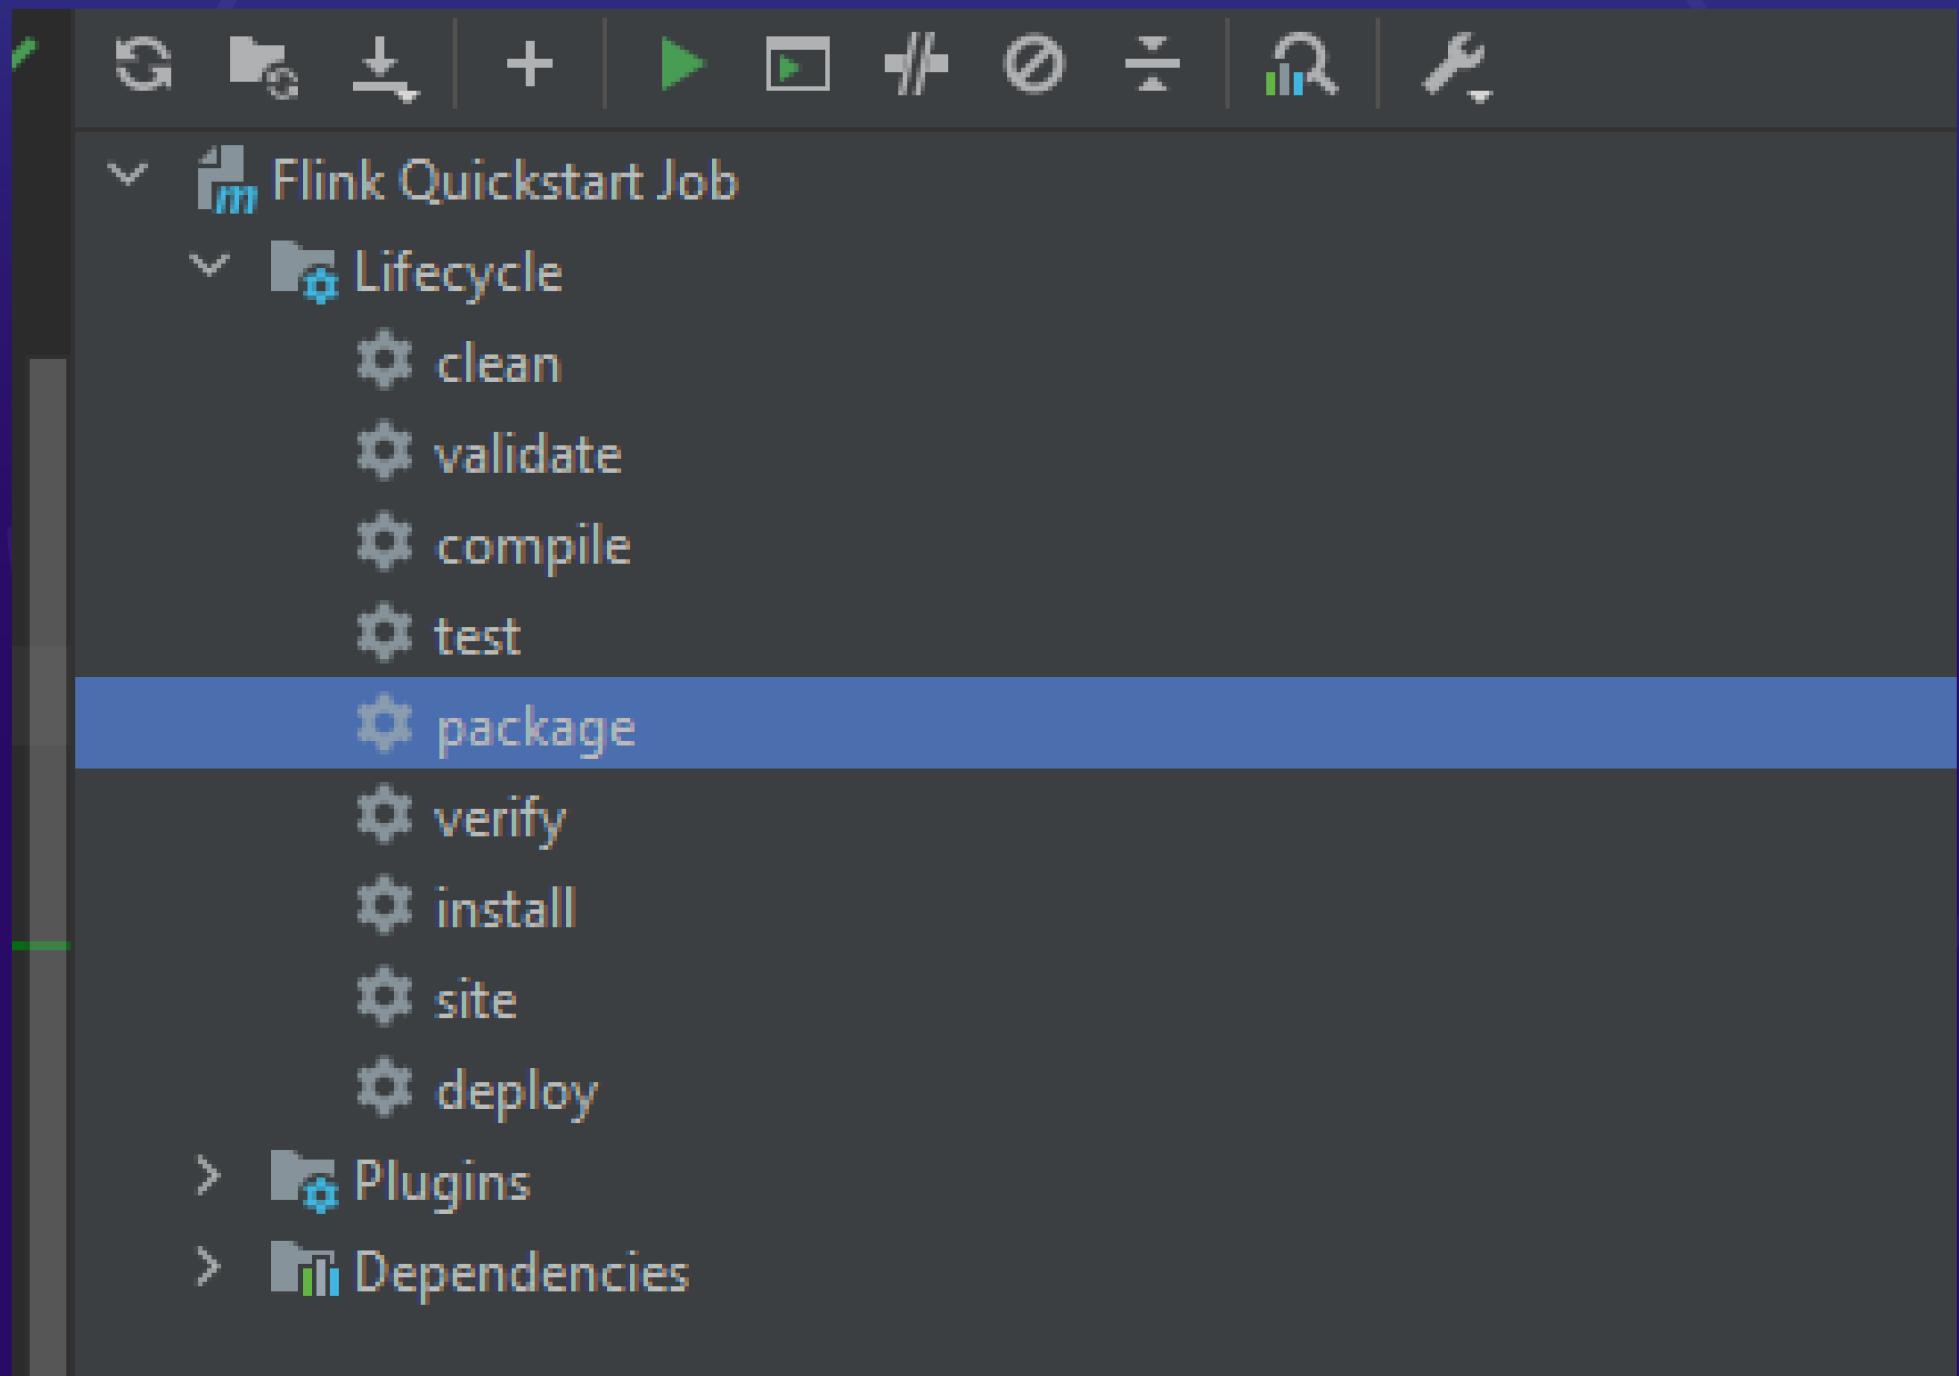
        LOGGER.info("Creating car data to sink into cassandraDB.");
        SingleOutputStreamOperator<Tuple5<String, String, String, String, String>> sinkLocationStream = sinkFilteredLocationStream.map(
            (MapFunction<Tuple5<String, Double, Double, Double, Double>, Tuple5<String, String, String, String, String>>) filteredData ->
                new Tuple5<>(filteredData.f0, Double.toString(filteredData.f1), Double.toString(filteredData.f2), Double.toString(filteredData.f3), Double.toString(filteredData.f4)))
            .returns(new TupleTypeInfo<>(TypeInformation.of(String.class),
                TypeInformation.of(String.class),
                TypeInformation.of(String.class),
                TypeInformation.of(String.class),
                TypeInformation.of(String.class)));

        sinkLocationStream.print();
        LOGGER.info("Open Cassandra connection and Sinking locations data into cassandraDB.");
        CassandraSink
            .addSink(sinkLocationStream)
        //...
        .setHost("cassandra")
        .setQuery("INSERT INTO locations_db.flinkgeolocation(user, minvalue, maxvalue, meanvalue, count) values (?, ?, ?, ?, ?);")
        .build();
    }
}
```



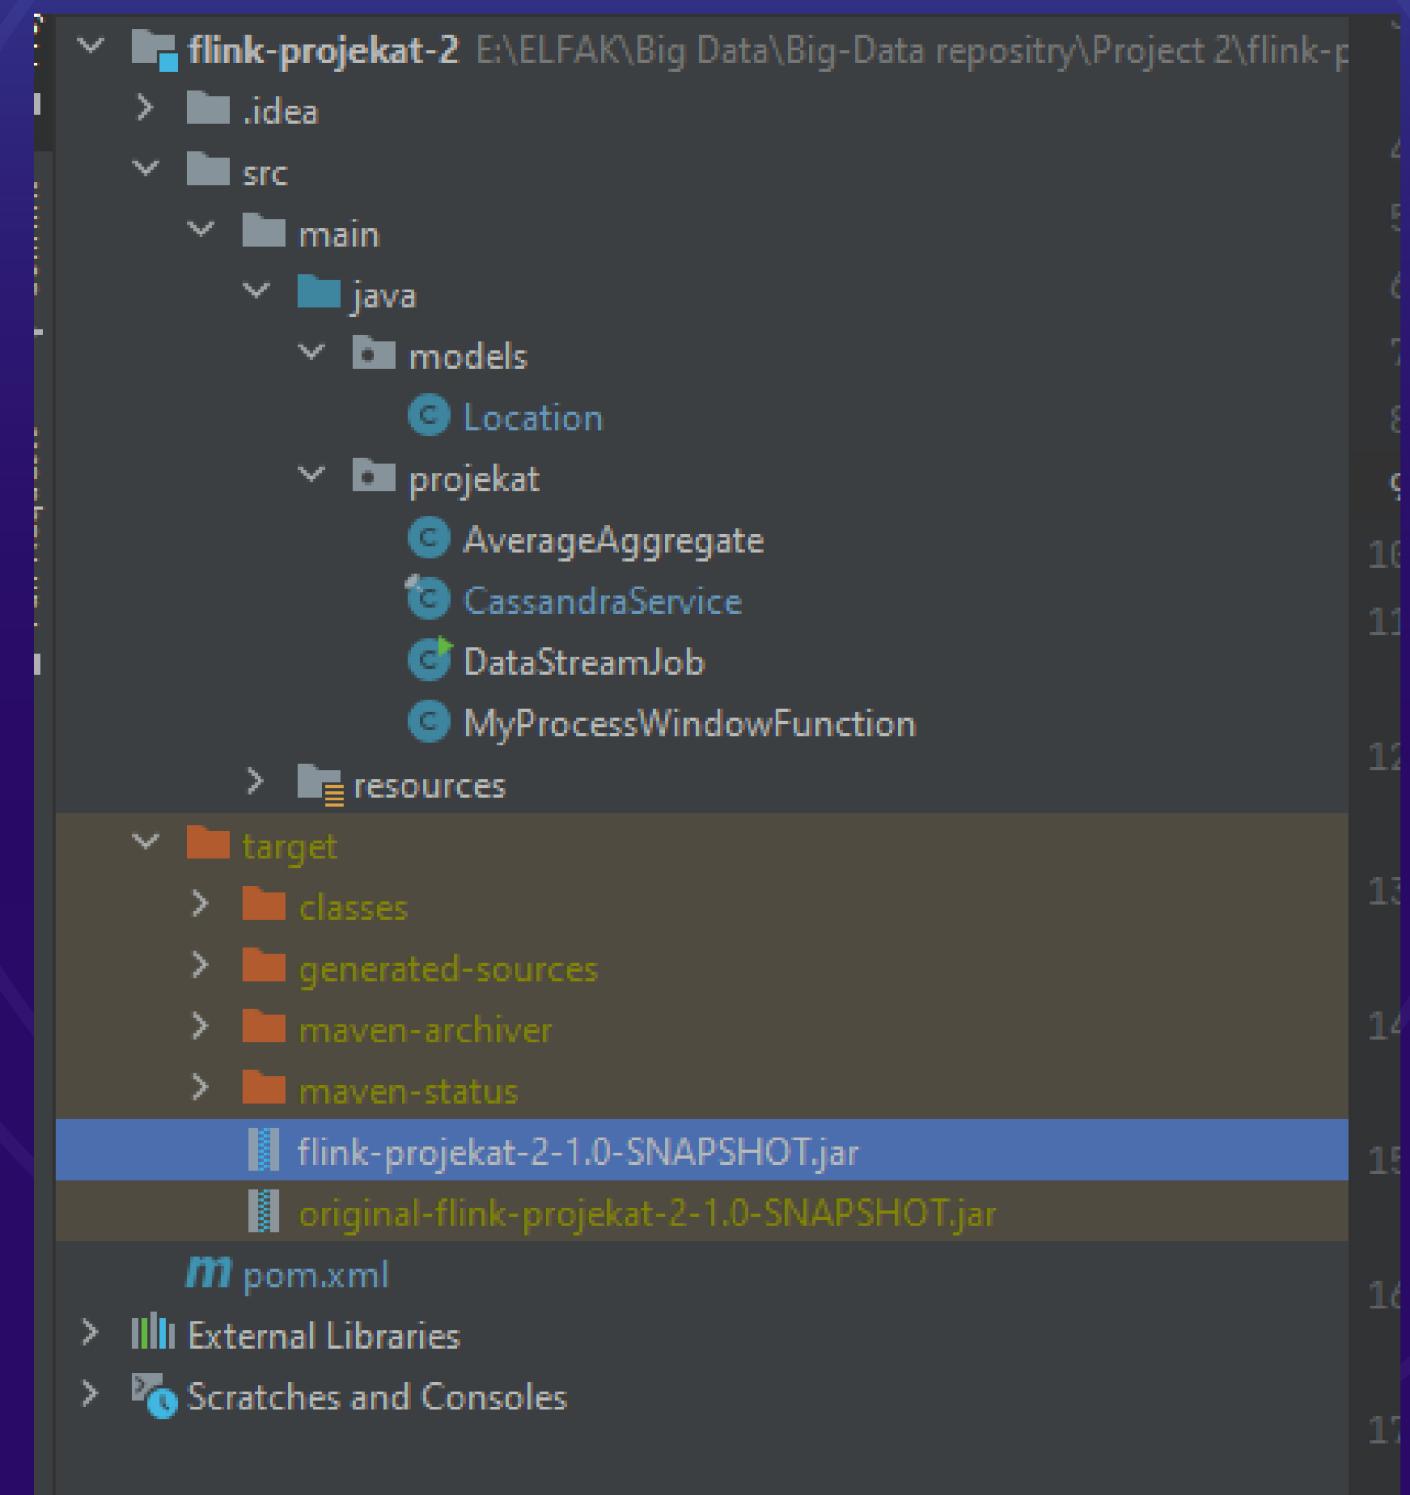
# Flink Streaming aplikacija

## Kreiranje .jar fajla uz pomoć Maven package



# Flink Streaming aplikacija

Dobija se SNAPSHOT.jar fajl



# Flink Streaming aplikacija

## Postavljanje .jar fajla na Flink kontejner

Uploaded Jars

Name	Upload Time	Entry Class	
flink-projekat-2-1.0-SNAPSHOT.jar	2023-04-04, 19:58:20	projekat.DataStreamJob	<a href="#">Delete</a>

[projekat.DataStreamJob](#)

[Parallelism](#)

[Program Arguments](#)

[Savepoint Path](#)

Allow Non Restored State

[Show Plan](#) [Submit](#)



# Flink Streaming aplikacija

## Pregled pokrenutog posla na Flink kontejner-u

Flink Streaming Job

Cancel Job

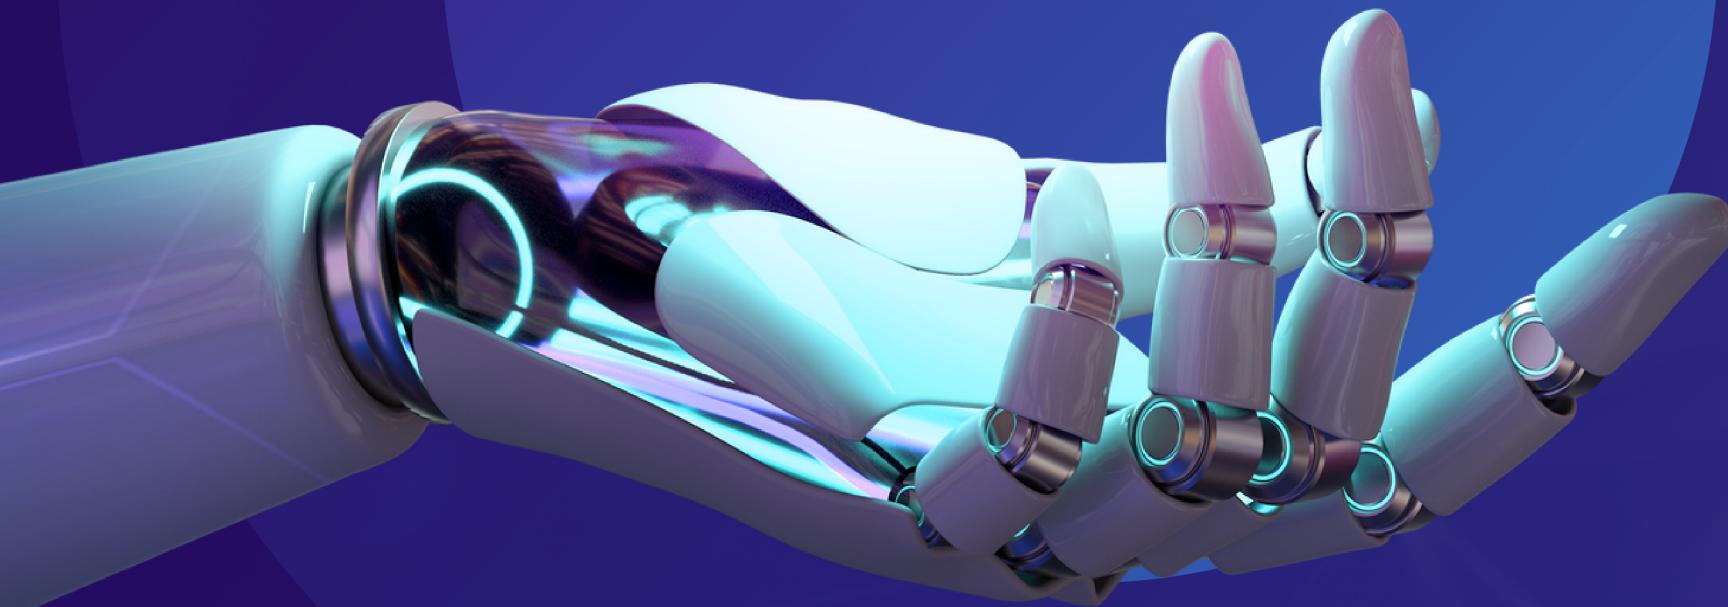
Job ID	8700c63a685b98e702b94b155168fb9b	Job State	RUNNING   3	Actions	Job Manager Log	
Start Time	2023-04-04 19:58:42	Duration	9s			

Overview   Exceptions   Timeline   Checkpoints   Configuration

```
graph LR; A["Source: Custom Source -> Map -> Filter  
Parallelism: 1  
Backpressured (max): 0%  
Busy (max): N/A"] -- HASH --> B["SlidingProcessingTimeWindows -> Map -> (Sink: Print to Std. Out, Sink: Cassandra Sink)  
Parallelism: 1  
Backpressured (max): 0%  
Busy (max): N/A"]; A -- HASH --> C["SlidingProcessingTimeWindows  
Parallelism: 1  
Backpressured (max): 0%  
Busy (max): N/A"]
```

Name	Status	Bytes Received	Records Received	Bytes Sent	Records Sent	Parallelism	Start Time	Duration	Tasks
Source: Custom Source -> Map -> Map -> Filter	RUNNING	0 B	0	0 B	0	1	2023-04-04 19:58:43	9s	1
SlidingProcessingTimeWindows -> Map -> (Sink: Print to Std. Out, Sink: Cassandra Sink)	RUNNING	0 B	0	0 B	0	1	2023-04-04 19:58:43	9s	1
SlidingProcessingTimeWindows	RUNNING	0 B	0	0 B	0	1	2023-04-04 19:58:43	9s	1

→



HVALA  
NA  
PAŽNJI



[ficax.ai@elfak.rs](mailto:ficax.ai@elfak.rs)

